

# 911 Calls Capstone Project - Solutions

For this project, I analyzed some 911 call data from [Kaggle](#). Being an exploratory analysis to better understand how emergency calls behave, where we can also answer where we should place facilities to better respond to these emergency calls.

The data contains the following fields:

- lat : String variable, Latitude
- lng: String variable, Longitude
- desc: String variable, Description of the Emergency Call
- zip: String variable, Zipcode
- title: String variable, Title
- timeStamp: String variable, YYYY-MM-DD HH:MM:SS
- twp: String variable, Township
- addr: String variable, Address
- e: String variable, Dummy variable (always 1)

However, I made some changes to the dataset to try to fix it with python, so I put some letters at the end of the values in the lat and lng columns, added some empty rows and columns.

---

## Setup

### Importing libraries that will be used

```
In [1]: import pandas as pd
import numpy as np
```

### Importing visualization libraries

```
In [2]: import matplotlib.pyplot as plt
import seaborn as sns
sns.set_style('whitegrid') # change the general style of the plots (just because I like
import io
from PIL import Image
from IPython import display
import cufflinks as cf
cf.go_offline()

# Displays output inline
%matplotlib inline
```

---

## Loading the data

### Loading the csv file as a dataframe called df

```
In [3]: df = pd.read_csv('911_edit.csv', sep=';') # This file is too big to put on Github, so if
```

```
df.head(3)
```

Out[3]:

	lat	lng	desc	zip	title	timeStamp	twp	addr	
0	40,2978759	-75,5812935nn	REINDEER CT & DEAD END; NEW HANOVER; Station ...	19525.0	EMS: BACK PAINS/INJURY	10/12/2015 17:40	NEW HANOVER	REINDEER CT & DEAD END	1
1	40,2580614nn	-75,2646799	BRIAR PATH & WHITEMARSH LN; HATFIELD TOWNSHIP...	19446.0	EMS: DIABETIC EMERGENCY	10/12/2015 17:40	HATFIELD TOWNSHIP	BRIAR PATH & WHITEMARSH LN	1
2	40,1211818	-75,3519752	HAWS AVE; NORRISTOWN; 2015-12-10 @ 14:39:21-St...	19401.0	Fire: GAS- ODOR/LEAK	10/12/2015 17:40	NORRISTOWN	HAWS AVE	1

There are some letter at the end of values columns in the lat and lng that need to be removed. Also there are three empty columns which aren't needed

## Pre-processing

In [4]:

```
# removing null collumns
df.dropna(axis=1, how="all", inplace=True) # parameter HOW because we want to remove colu
print(df.columns)
```

```
Index(['lat', 'lng', 'desc', 'zip', 'title', 'timeStamp', 'twp', 'addr', 'e'], dtype='object')
```

In [5]:

```
# seeing if there are empty rows
df['lat'].isnull().value_counts()
```

Out[5]:

```
False    99492
True       17
Name: lat, dtype: int64
```

In [6]:

```
# removing null rows
df.dropna(axis=0, how='all', inplace=True)
df['lat'].isnull().value_counts()
```

Out[6]:

```
False    99492
Name: lat, dtype: int64
```

## Function

In [7]:

```
# fix the letters in the lat and lng columns
def remove_nn(value):

    if type(value) is str:
        return value.replace('n', '')
    else:
        return value

# test
remove_nn("406030nn")
```

Out[7]:

```
'406030'
```

```
In [8]: # apply in the columns
df['lat'] = df['lat'].apply(remove_nn)
df['lng'] = df['lng'].apply(remove_nn)
df.head(5)
```

	lat	lng	desc	zip	title	timeStamp	twp	addr	e
0	40,2978759	-75,5812935	REINDEER CT & DEAD END; NEW HANOVER; Station ...	19525.0	EMS: BACK PAINS/INJURY	10/12/2015 17:40	NEW HANOVER	REINDEER CT & DEAD END	1.0
1	40,2580614	-75,2646799	BRIAR PATH & WHITEMARSH LN; HATFIELD TOWNSHIP...	19446.0	EMS: DIABETIC EMERGENCY	10/12/2015 17:40	HATFIELD TOWNSHIP	BRIAR PATH & WHITEMARSH LN	1.0
2	40,1211818	-75,3519752	HAWS AVE; NORRISTOWN; 2015-12-10 @ 14:39:21-St...	19401.0	Fire: GAS-ODOR/LEAK	10/12/2015 17:40	NORRISTOWN	HAWS AVE	1.0
3	40,116153	-75,343513	AIRY ST & SWEDE ST; NORRISTOWN; Station 308A;...	19401.0	EMS: CARDIAC EMERGENCY	10/12/2015 17:40	NORRISTOWN	AIRY ST & SWEDE ST	1.0
4	40,251492	-75,6033497	CHERRYWOOD CT & DEAD END; LOWER POTTSGROVE; S...	NaN	EMS: DIZZINESS	10/12/2015 17:40	LOWER POTTSGROVE	CHERRYWOOD CT & DEAD END	1.0

```
In [9]: # The changes that I made broke the timeStamp column, so I will replace the column of th
df_or = pd.read_csv("911.csv")
df_or.head()
```

	lat	lng	desc	zip	title	timeStamp	twp	addr	e
0	40.297876	-75.581294	REINDEER CT & DEAD END; NEW HANOVER; Station ...	19525.0	EMS: BACK PAINS/INJURY	2015-12-10 17:10:52	NEW HANOVER	REINDEER CT & DEAD END	1
1	40.258061	-75.264680	BRIAR PATH & WHITEMARSH LN; HATFIELD TOWNSHIP...	19446.0	EMS: DIABETIC EMERGENCY	2015-12-10 17:29:21	HATFIELD TOWNSHIP	BRIAR PATH & WHITEMARSH LN	1
2	40.121182	-75.351975	HAWS AVE; NORRISTOWN; 2015-12-10 @ 14:39:21-St...	19401.0	Fire: GAS-ODOR/LEAK	2015-12-10 14:39:21	NORRISTOWN	HAWS AVE	1
3	40.116153	-75.343513	AIRY ST & SWEDE ST; NORRISTOWN; Station 308A;...	19401.0	EMS: CARDIAC EMERGENCY	2015-12-10 16:47:36	NORRISTOWN	AIRY ST & SWEDE ST	1
4	40.251492	-75.603350	CHERRYWOOD CT & DEAD END; LOWER POTTSGROVE; S...	NaN	EMS: DIZZINESS	2015-12-10 16:56:52	LOWER POTTSGROVE	CHERRYWOOD CT & DEAD END	1

```
In [10]: df['timeStamp'] = df_or['timeStamp']
df.head()
```

```
Out[10]:
```

	lat	lng	desc	zip	title	timeStamp	twp	addr	e
0	40,2978759	-75,5812935	REINDEER CT & DEAD END; NEW HANOVER; Station ...	19525.0	EMS: BACK PAINS/INJURY	2015-12-10 17:10:52	NEW HANOVER	REINDEER CT & DEAD END	1.0
1	40,2580614	-75,2646799	BRIAR PATH & WHITEMARSH LN; HATFIELD TOWNSHIP...	19446.0	EMS: DIABETIC EMERGENCY	2015-12-10 17:29:21	HATFIELD TOWNSHIP	BRIAR PATH & WHITEMARSH LN	1.0
2	40,1211818	-75,3519752	HAWS AVE; NORRISTOWN; 2015-12-10 @ 14:39:21-St...	19401.0	Fire: GAS- ODOR/LEAK	2015-12-10 14:39:21	NORRISTOWN	HAWS AVE	1.0
3	40,116153	-75,343513	AIRY ST & SWEDE ST; NORRISTOWN; Station 308A;...	19401.0	EMS: CARDIAC EMERGENCY	2015-12-10 16:47:36	NORRISTOWN	AIRY ST & SWEDE ST	1.0
4	40,251492	-75,6033497	CHERRYWOOD CT & DEAD END; LOWER POTTSGROVE; S...	NaN	EMS: DIZZINESS	2015-12-10 16:56:52	LOWER POTTSGROVE	CHERRYWOOD CT & DEAD END	1.0

### Check the info() of the df

```
In [11]: df.info()

<class 'pandas.core.frame.DataFrame'>
Int64Index: 99492 entries, 0 to 99508
Data columns (total 9 columns):
#   Column      Non-Null Count  Dtype
---  -
0   lat         99492 non-null  object
1   lng         99492 non-null  object
2   desc        99492 non-null  object
3   zip         86637 non-null  float64
4   title       99492 non-null  object
5   timeStamp   99492 non-null  object
6   twp         99449 non-null  object
7   addr        98973 non-null  object
8   e           99492 non-null  float64
dtypes: float64(2), object(7)
memory usage: 7.6+ MB
```

```
In [12]: # Converting the lat and lng to float
""" obs: My excel is in Portuguese so when I changed the values in the first two columns
the decimal separator changed. So I needed to switch to """

df['lat'] = df['lat'].str.replace(',','.' ).astype(float)
df['lng'] = df['lng'].str.replace(',','.' ).astype(float)

# Converting "e" column to int64 (just to be the same as the original dataset)
df['e'] = df['e'].astype('int64')
```

```
In [13]: # See the type of "timeStamp" column
type(df['timeStamp'].iloc[0])
```

Out[13]: str

```
In [14]: # We can't work with data in str type, so it's necessary to convert the timeStamp column  
df['timeStamp'] = pd.to_datetime(df['timeStamp'])
```

```
In [15]: type(df['timeStamp'].iloc[0]) # now we can manipulate dates and times  
  
# Work with timeStamp becomes possible grab specific attributes from a Datetime object b
```

Out[15]: pandas.\_libs.tslibs.timestamps.Timestamp

```
In [16]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
Int64Index: 99492 entries, 0 to 99508  
Data columns (total 9 columns):  
#   Column      Non-Null Count  Dtype  
---  ---  
0   lat         99492 non-null   float64  
1   lng         99492 non-null   float64  
2   desc        99492 non-null   object  
3   zip         86637 non-null   float64  
4   title       99492 non-null   object  
5   timeStamp   99492 non-null   datetime64[ns]  
6   twp         99449 non-null   object  
7   addr        98973 non-null   object  
8   e           99492 non-null   int64  
dtypes: datetime64[ns](1), float64(3), int64(1), object(4)  
memory usage: 7.6+ MB
```

```
In [17]: # Removing outliers  
# We have 3 latitudes that are very different from the other values, indicating that the  
df.lat.sort_values()
```

```
Out[17]: 67120    30.333596  
25726    32.387090  
57273    39.745533  
88516    39.745533  
75249    39.956497  
...  
13604    40.584929  
57577    40.584929  
71621    40.584929  
7761     40.584929  
8378     41.167156  
Name: lat, Length: 99492, dtype: float64
```

```
In [18]: df.lng.sort_values()
```

```
Out[18]: 67120    -95.595595  
25726    -86.276106  
57273    -84.395256  
88516    -84.395256  
63388    -77.686813  
...  
48628    -75.012379  
11295    -75.008877  
52019    -75.008041  
1528     -75.000755  
52063    -74.995041  
Name: lng, Length: 99492, dtype: float64
```

```
In [19]: df.drop([67120,25726,8378], inplace=True) # Removing by index that we can see above
```

Thus, the pre-processing step ends here

---

## Analytics

### Find the top 5 zipcodes for 911 calls

```
In [20]: df['zip'].value_counts().head(5)
```

```
Out[20]: 19401.0    6979
19464.0    6643
19403.0    4854
19446.0    4748
19406.0    3174
Name: zip, dtype: int64
```

### Top 5 townships (twp) for 911 calls

In Brazil, we doesn't have this kind of division so it was a little strange for me. I assumed it's like cities (many quotes hahah)

```
In [21]: df['twp'].value_counts().head(5)
```

```
Out[21]: LOWER MERION    8443
ABINGTON    5977
NORRISTOWN   5890
UPPER MERION 5227
CHELTENHAM   4575
Name: twp, dtype: int64
```

### Unique title codes

To see how many kinds of emergency kinds are in the dataset

```
In [22]: df['title'].nunique()
```

```
Out[22]: 110
```

## Creating new features

### In the titles column there are "Reasons/Departments" specified before the title code

```
In [23]: # Create a new column called "Reason" that contains only the title code **
df['Reason'] = df['title'].apply(lambda title: title.split(':')[0])
```

```
In [24]: reason_count = df['Reason'].value_counts()
reason_count_p = df['Reason'].value_counts(normalize=True)
pd.concat([reason_count, round(reason_count_p,2)], axis=1, keys=['count', 'percentage'])
```

```
Out[24]:
```

	count	percentage
<b>EMS</b>	48875	0.49
<b>Traffic</b>	35694	0.36
<b>Fire</b>	14920	0.15


We can see that the dataset has three categories EMS, Fire, and Traffic.

Also, the most common emergency is EMS (almost 50% of the cases) which is a emergency services that provide urgent pre-hospital treatment and stabilisation for serious illness and injuries and transport to definitive care. It is important to note that this classification can be confused with "Traffic" and vice versa.

### Using seaborn to create a countplot of 911 calls by Reason.

```
In [25]: plt.figure(figsize=(4,3))
sns.countplot(x='Reason',data=df,palette='viridis')
```

```
Out[25]: <AxesSubplot:xlabel='Reason', ylabel='count'>
```



To focus on time information, I created 3 new columns called Hour, Month, and Day of Week to understand the distribution and whether there are any patterns

```
In [26]: # using the lambda function as it is a simple and punctual function

df['Hour'] = df['timeStamp'].apply(lambda time: time.hour)
df['Month'] = df['timeStamp'].apply(lambda time: time.month)
df['Year'] = df['timeStamp'].apply(lambda time: time.year)
df['Day of Week'] = df['timeStamp'].apply(lambda time: time.dayofweek)
df.head(2)
```

```
Out[26]:
```

	lat	lng	desc	zip	title	timeStamp	twp	addr	e	Reason
0	40.297876	-75.581294	REINDEER CT & DEAD END; NEW HANOVER; Station ...	19525.0	EMS: BACK PAINS/INJURY	2015-12-10 17:10:52	NEW HANOVER	REINDEER CT & DEAD END	1	EMS
1	40.258061	-75.264680	BRIAR PATH & WHITEMARSH LN; HATFIELD TOWNSHIP...	19446.0	EMS: DIABETIC EMERGENCY	2015-12-10 17:29:21	HATFIELD TOWNSHIP	BRIAR PATH & WHITEMARSH LN	1	EMS

The Day of Week is an integer 0-6. So using the .map() with a dictionary to map the actual string names to the day of the week:

```
In [27]: dmap = {0:'Mon',1:'Tue',2:'Wed',3:'Thu',4:'Fri',5:'Sat',6:'Sun'}
```

```
In [28]: df['Day of Week'] = df['Day of Week'].map(dmap)
df.head(2)
```

Out[28]:

	lat	lng	desc	zip	title	timeStamp	twp	addr	e	Reason
0	40.297876	-75.581294	REINDEER CT & DEAD END; NEW HANOVER; Station ...	19525.0	EMS: BACK PAINS/INJURY	2015-12-10 17:10:52	NEW HANOVER	REINDEER CT & DEAD END	1	EMS
1	40.258061	-75.264680	BRIAR PATH & WHITEMARSH LN; HATFIELD TOWNSHIP...	19446.0	EMS: DIABETIC EMERGENCY	2015-12-10 17:29:21	HATFIELD TOWNSHIP	BRIAR PATH & WHITEMARSH LN	1	EMS

```
In [29]: print(f"Min. Date: {df['timeStamp'].min()}")
print(f"Max. Date: {df['timeStamp'].max()}")
```

Min. Date: 2015-12-10 14:39:21  
Max. Date: 2016-08-25 19:10:15

**Using seaborn to create a countplot of the Day of Week column with the hue based off of the Reason column.**


```
In [30]: sns.countplot(x='Day of Week',data=df,hue='Reason',palette='viridis', )
print(df.groupby('Day of Week')['Reason'].count().sort_values(ascending=False))

# To relocate the legend
plt.legend(bbox_to_anchor=(1.05, 1), loc=2, borderaxespad=0.)
```

Day of Week  
Tue 15057  
Wed 15043  
Fri 14763  
Thu 14679  
Mon 14625  
Sat 13236  
Sun 12086  
Name: Reason, dtype: int64  
<matplotlib.legend.Legend at 0x1da43bf39d0>

Out[30]:






So Tuesday was the Day of Week which had more Emergency Calls, followed by Wednesday. Other point is that Fire was almost constant. When we look to Traffic, we can see an interesting pattern, where Sunday had the smaller frequency which could be explained by the characteristic about the day, because people don't go out a lot in Sunday.

**Now do the same for Month:**

```
In [31]: plt.figure(figsize=(4,3))
sns.countplot(x='Month', data=df, hue='Reason', palette='viridis')

# To relocate the legend
plt.legend(bbox_to_anchor=(1.05, 1), loc=2, borderaxespad=0.)
```

Out[31]: <matplotlib.legend.Legend at 0x1da5a7ba100>



January was the month with the most emergency calls, followed by July and June.

For some reason, it is missing some months -> 9,10, and 11 are not there

**We can fill in this information by plotting the information in another way, possibly a simple line plot that fills in the missing months**

```
In [32]: byMonth = df.groupby('Month').count()  
byMonth.head()
```


```
Out[32]:
```

	lat	lng	desc	zip	title	timeStamp	twp	addr	e	Reason	Hour	Year	Day of Week
Month													
1	13095	13095	13095	11421	13095	13095	13094	12986	13095	13095	13095	13095	13095
2	11395	11395	11395	9861	11395	11395	11393	11324	11395	11395	11395	11395	11395
3	11059	11059	11059	9737	11059	11059	11050	11018	11059	11059	11059	11059	11059
4	11287	11287	11287	9856	11287	11287	11284	11244	11287	11287	11287	11287	11287
5	11374	11374	11374	9914	11374	11374	11371	11330	11374	11374	11374	11374	11374

**Now we can create a simple plot off of the dataframe indicating the count of calls per month.**

```
In [33]: plt.figure(figsize=(10,4))  
byMonth['Reason'].plot()
```

```
Out[33]: <AxesSubplot: xlabel='Month'>
```




Although we do not have data for months 9, 10 and 11, it is likely that it followed the trend shown in the graph above.

**Using seaborn lmpplot() to create a linear fit on the number of calls per month, i.e., see the regression line**

```
In [34]: sns.lmplot(x='Month',y='twp',data=byMonth.reset_index(),height=3, aspect=2)  
# ps: We can use any column to do the plot, because is a counting in y axis and the valu
```

```
Out[34]: <seaborn.axisgrid.FacetGrid at 0x1da5791ad00>
```



So, as I said before and looking at the regression line above, we were able to at least get an idea of how the data could be in the months that had no data

### Creating a new column called 'Date' that contains the date from the timeStam column for to see the distribution per date

```
In [35]: df['Date'] = df['timeStamp'].apply(lambda t: t.date())
df['Date'].head(1)
```

```
Out[35]: 0    2015-12-10
Name: Date, dtype: object
```

```
In [36]: # mean of calls per date
print(f"Mean: {df.groupby('Date')['Reason'].count().mean()}")

# median of calls per date
print(f"Median: {df.groupby('Date')['Reason'].count().median()}")
```

```
Mean: 382.65
Median: 381.5
```

Both calculations were made to understand the data, as we have that the mean is more susceptible to extreme values, so it is normally used when the data are arranged homogeneously. The median is less influenced by very high or very low values. Despite the two values being close, I adopted the median as the standard for the other calculations in order to avoid problems with days that may have connection peaks.

```
In [37]: # median of calls per reason
ems_median = df[df['Reason'] == 'EMS'].groupby('Date')['Reason'].count().median()
traffic_median = df[df['Reason'] == 'Traffic'].groupby('Date')['Reason'].count().median()
fire_median = df[df['Reason'] == 'Fire'].groupby('Date')['Reason'].count().median()

print('Median by Date:\n')
print(f"EMS: {ems_median}")
print(f"Traffic: {traffic_median}")
print(f"Fire: {fire_median}")
```

```
Median by Date:
```

```
EMS: 189.0
Traffic: 134.0
Fire: 56.0
```


```
In [38]: # Top 3 dates with the most numbers of calls
df.groupby('Date')['Reason'].count().sort_values(ascending=False).head(3)
```

```
Out[38]: Date
2016-01-23    887
2016-02-24    673
2016-01-24    657
Name: Reason, dtype: int64
```

### Groupby this Date column with the count() aggregate and create a plot of counts of 911 calls

I used cufflinks because it's more interactive and helps a lot with graph analysis. However, to be able to see it here on GitHub, I took a print of the graph generated by the commented code, that's why it has Image(filename)

```
In [39]: #df.groupby('Date')['Reason'].count().iplot(theme="white", color="purple")
display.Image('all.png')
```



```
In [40]: print(f"\033[1mSo Jan 23, 2016 was {round(((931-384)/384)*100)}% above the median for em
So Jan 23, 2016 was 142% above the median for emergency calls
```

We can see that on Jan 23, 2016 there was a spike in emergency calls (931 calls). Researching I found that on the 22nd to the 24th there was an intense blizzard in the northeastern of the USA.

(<https://www.weather.gov/media/lwx/Bliz2016.pdf>)


### Now recreating this plot but create 3 separate plots with each plot representing a Reason for the 911 call

If exists some correlation with the spike in the Jan 23, 2016 with the blizzard probably emergency calls by "Traffic" will be bigger than other Reason, because in blizzards the streets/roads are more dangerous to drive.

Blizzards combine the limited visibility of fog with the slippery roads you can expect from ice and snow. Driving in a blizzard makes it incredibly difficult to see what is going on around you, be seen by other road users and maintain control of your vehicle. (<https://www.epermittest.com/drivers-education/driving-blizzard>)


```
In [41]: #df[df['Reason']=='Traffic'].groupby('Date').count()['twp'].iplot(title="Traffic", colors
display.Image('traffic_date.png')
```

Out[41]: Traffic




```
In [42]: #df[df['Reason']=='Fire'].groupby('Date').count()['twp'].iplot(title="Fire",colors="red")
display.Image('fire_date.png')
```

Out[42]: Fire



```
In [43]: #df[df['Reason']=='EMS'].groupby('Date').count()['twp'].iplot(title="EMS",theme='white')
display.Image('ems_date.png')
```

Out[43]:



## Trying show all 3 graph in just 1

I spent a long time thinking about how to do it and I don't know if the way I did it was the best hahah. So if you know a simple way or another way to do this let me know.

```
In [44]: multi_index_df = df.groupby(by=['Date', 'Reason']).count()['twp']
multi_index_df
```

```
Out[44]:
```

Date	Reason	
2015-12-10	EMS	51
	Fire	12
	Traffic	40
2015-12-11	EMS	183
	Fire	69
2016-08-24	Fire	53
	Traffic	136
	2016-08-25	EMS
	Fire	49
	Traffic	89

Name: twp, Length: 780, dtype: int64


```
In [45]: df_graphics = multi_index_df.unstack()
df_graphics.head(3)
```

```
Out[45]:
```

	Reason	EMS	Fire	Traffic
Date				
2015-12-10		51	12	40
2015-12-11		183	69	138
2015-12-12		189	68	145

```
In [46]: #df_graphics.iplot(theme='white')
display.Image('allinone.png')
```

```
Out[46]:
```



```
In [47]: # % of Traffic in all emergency calls in 23 Jan,2016
print(f"{{ (559/913) :.0%}}")

# Traffic above median per date
print(f"{{ ((559-traffic_median)/traffic_median) :.0%}}")

# Fire above median
print(f"{{ ((158-fire_median)/fire_median) :.0%}}")

61%
317%
182%
```

**As we can see, the most common type of emergency call on Jan 23, 2016 was "Traffic" (61% of all calls) and that day it was above the daily median in 308%. Another point is that we have a peak in calls of the "Fire" type, going from a daily median of 55 to 158 (an increase of 187%)**

"Traffic" has more spikes and more fluctuations compared to others.




Finally, we can see that at the beginning of April we have a "stabilization" of the values, that is, they are closer to the medians of each type of emergency call and maintain a more constant behavior

---

Plotly and Cufflinks are great for data visualization. It's possible to separate the graph as a "subplot" in matplotlib just by adding the parameter (subplots=True), in addition, we can customize it in several ways due to the various parameters that lplot has

```
In [48]: #df_graphics.iplot(subplots=True, theme='white')
display.Image('threegraphs.png')
```

Out[48]:



## Heatmaps

### Day of Week vs Month

Now creating heatmaps with seaborn and our data. First is needed to restructure the dataframe

A heatmap is a graphical representation of data where each value of a matrix is represented as a color, this makes it easier to understand the information understanding of the information

```
In [49]: dayMonth = df.groupby(by=['Day of Week', 'Month']).count()['Reason'].unstack()
          dayMonth
```


```
Out[49]:
```

	Month	1	2	3	4	5	6	7	8	12
<b>Day of Week</b>										
<b>Fri</b>	1950	1564	1521	1949	1725	1646	2042	1308	1058	
<b>Mon</b>	1718	1959	1534	1594	1776	1609	1685	1499	1251	
<b>Sat</b>	2252	1434	1264	1736	1438	1376	1691	1085	960	
<b>Sun</b>	1947	1223	1097	1486	1413	1329	1667	1020	904	
<b>Thu</b>	1584	1582	1883	1592	1585	2055	1642	1505	1251	
<b>Tue</b>	1961	1743	1871	1421	1904	1671	1650	1602	1234	
<b>Wed</b>	1683	1890	1889	1509	1533	2045	1711	1536	1247	

```
In [50]: plt.figure(figsize=(12,4))
          sns.heatmap(dayMonth,cmap='coolwarm')
```

```
Out[50]: <AxesSubplot:xlabel='Month', ylabel='Day of Week'>
```





When we look at the day of the week by month, Saturdays in January stand out with a high concentration of the number of cases

## Day of Week vs Hour

Now repeating these same plots and operations, for a DataFrame that shows the Day of Week vs Hour

```
In [51]: dayHour = df.groupby(by=['Day of Week', 'Hour']).count()['Reason'].unstack()
dayHour
```

```
Out[51]:
```

	Hour	0	1	2	3	4	5	6	7	8	9	...	14	15	16	17	18	19	20	21	22
<b>Day of Week</b>																					
<b>Fri</b>		271	227	187	184	201	209	368	635	723	749	...	927	971	1036	972	793	720	636	561	518
<b>Mon</b>		280	213	213	182	213	267	404	679	828	795	...	860	939	964	994	878	732	611	471	463
<b>Sat</b>		369	293	257	254	224	243	262	396	479	630	...	771	800	813	765	767	703	609	561	498
<b>Sun</b>		362	316	280	266	233	246	301	408	496	620	...	692	671	668	717	673	648	519	464	388
<b>Thu</b>		283	204	237	163	179	223	386	607	789	848	...	901	989	935	1034	815	685	610	538	409
<b>Tue</b>		267	237	184	167	216	252	421	674	884	851	...	948	941	1014	1019	867	722	642	560	434
<b>Wed</b>		241	221	179	215	154	253	433	728	858	807	...	927	901	1010	1062	882	686	675	564	494

7 rows × 24 columns

```
In [52]: # top 4 highest occurrences by day and hour
dayHour.unstack().nlargest(4)
```


```
Out[52]:
```

Hour	Day of Week	
17	Wed	1062
16	Fri	1036
17	Thu	1034
	Tue	1019

dtype: int64

```
In [53]: plt.figure(figsize=(12,4))
sns.heatmap(dayHour,cmap='coolwarm')
```

```
Out[53]: <AxesSubplot: xlabel='Hour', ylabel='Day of Week'>
```



Here we can see some interesting things.

- From 8am to 6pm are the times that have the most call records, this can be explained by the fact that it is business hours, that is, it is precisely when there are more people awake and also doing things, so the probability of an emergency happening increases due to of the amount of people.
- We can highlight that the times with the highest number of occurrences are at 4 pm and 5 pm. On the other hand, the times with the lowest number of occurrences are in the early morning between 2 am and 3 am.

## Reason vs Hour

**Now repeating these same plots and operations, for a DataFrame that shows the Reason vs Hour**

```
In [54]: hourReason = df.groupby(by=['Reason', 'Hour']).count()['Date'].unstack()
```

```
In [55]: hourReason.unstack().nlargest(7)
```


```
Out[55]:
```

Hour	Reason	
17	EMS	3163
16	EMS	3061
15	EMS	3000
14	EMS	2988
12	EMS	2978
13	EMS	2914
18	EMS	2739

dtype: int64

```
In [56]: plt.figure(figsize=(12,4))  
sns.heatmap(hourReason,cmap='coolwarm')
```

```
Out[56]: <AxesSubplot: xlabel='Hour', ylabel='Reason'>
```



Here we can see how occurrences are distributed by time and type. Thus, it is concluded that the values highlighted above (4 pm and 5 pm) come from an emergency call of the 'Traffic' type, which makes sense, since it is the rush hour when there are more people in traffic and who are more tired because of the work routine, so there is a greater probability of accidents of this type.

On the other hand, we can observe that the EMS are more constant occurrences during the day, mainly in the middle of the morning and beginning of the afternoon (I didn't think of any hypothesis for this). EMS as said has a median occurrence higher than the others, so although "Traffic" has much higher peaks, it does not have a median value greater than "EMS"

Finally, we have that occurrences of the "Fire" type are more punctual, so we don't have such outliers, but like the others, they occur more during business hours.

**Abstract: The more people doing day-to-day activities, the greater the chances of having an accident. Be careful :)**

## Maps

**I thought it was cool to show a real map of the city to see where the calls come from**

If you want to learn more about this kind of plot: <https://www.kaggle.com/code/alexisbcook/interactive-maps#The-data>

```
In [57]: # Importing libraries
import folium
import math
from folium import Choropleth, Circle, Marker
from folium.plugins import HeatMap, MarkerCluster

In [58]: # Separating by reason. I wanted to see the case highlighted, that is, 4 p.m. to 5 p.m.
map_traffic = df[(df['Reason'] == 'Traffic') & (df['Hour'].isin(range(16,18)))]

In [59]: m = folium.Map(location=[40.243821, -75.649367], tiles='cartodbpositron', zoom_start=8)

# Add points to the map
"""mc = MarkerCluster()
for idx, row in map_traffic.iterrows():
    if not math.isnan(row['lng']) and not math.isnan(row['lat']):
```

```

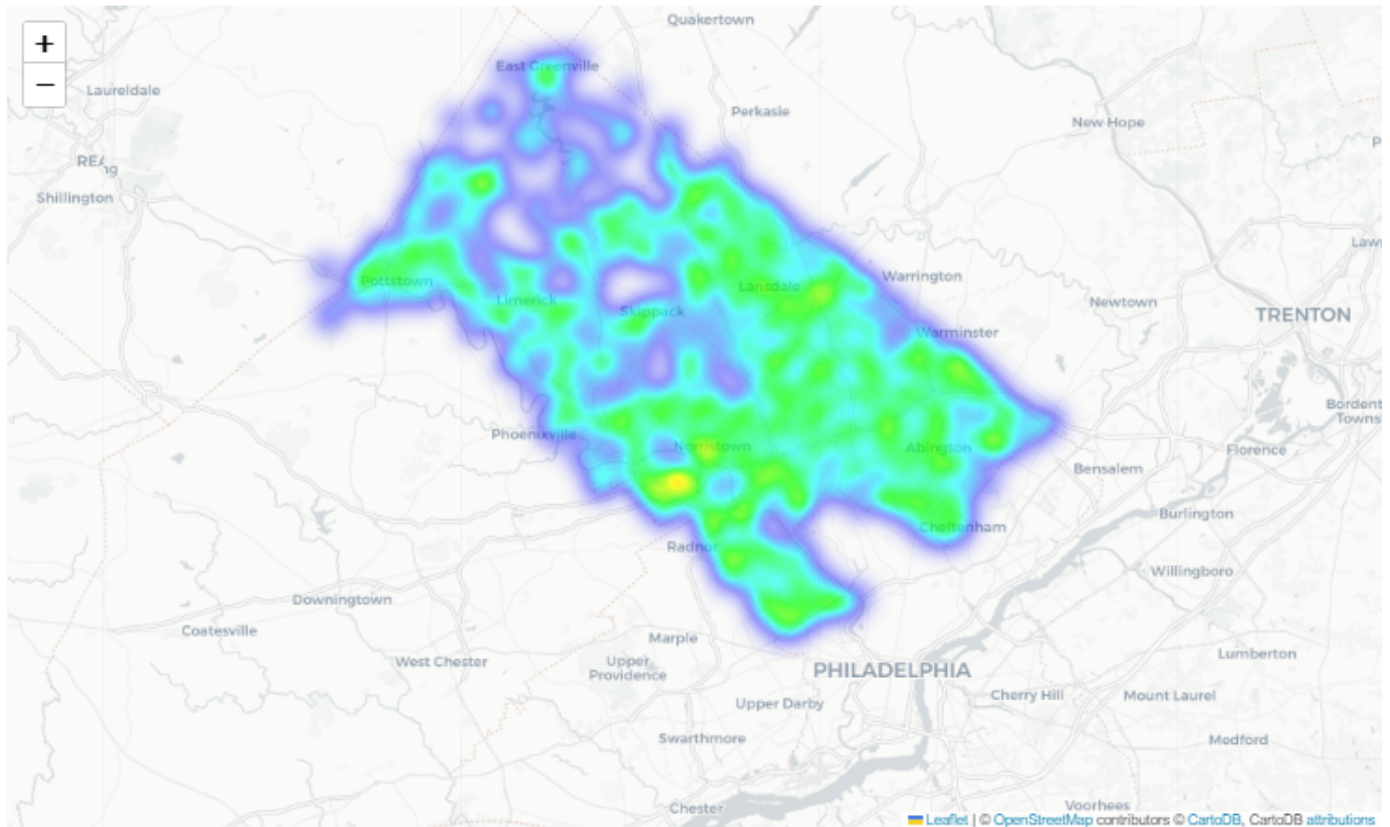
mc.add_child(Marker([row['lat'], row['lng']]))
m.add_child(mc)"""

# Add a heatmap to the base map
HeatMap(data=map_traffic[['lat', 'lng']], radius=10).add_to(m)

#Display the map
display.Image(filename='heatmap_traffic.png') #m

```

Out[59]:



Traffic accidents, as expected, occur more in centers and on major highways, mainly at intersections.

## Blizzard (22-24 Jan, 2016)

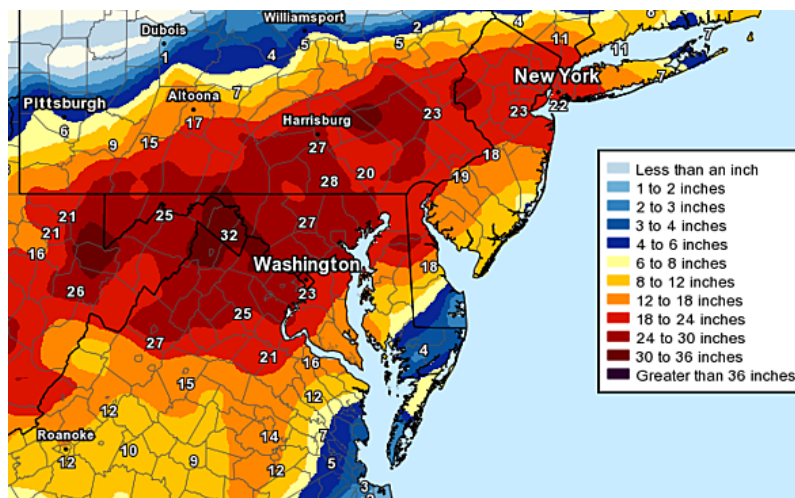
"The Blizzard of 2016 flooded coastal communities and piled up over 40 inches of snow, with incredible drifts. It set new snowstorm records at New York City and Baltimore."

```

In [60]: from IPython import display
display.Image("https://cms.accuweather.com/wp-content/uploads/2016/01/590x366_01252031_s

```

Out[60]:



```

In [61]: blizzard_traffic = df[(df['Reason'] == 'Traffic') & (df['timeStamp'] >= '2016-01-22') &

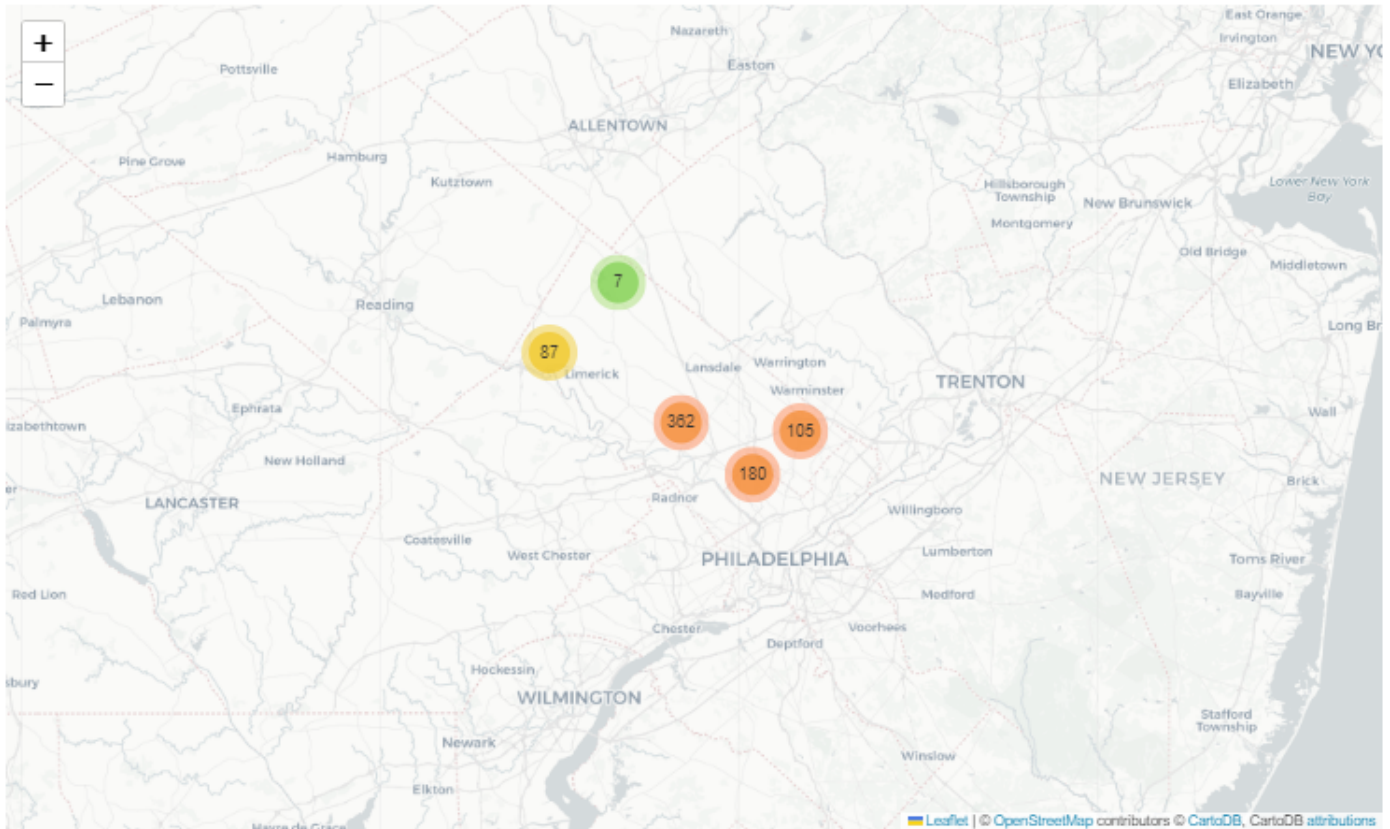
In [62]: # Create a base map
m_2 = folium.Map(location=[40.243821, -75.649367], tiles='cartodbpositron', zoom_start=8

# Add points to the map
mcd = MarkerCluster()
for idx, row in blizzard_traffic.iterrows():
    if not math.isnan(row['lng']) and not math.isnan(row['lat']):
        mcd.add_child(Marker([row['lat'], row['lng']]))
m_2.add_child(mcd)

# Display the map
display.Image('blizzard_traffic.png') #m_2

```

Out[62]:



Finally, I found interesting to see the EMS, since it is the type of emergency that has the highest frequency, so I limited it to the times with the highest concentration of cases, as seen earlier. The goal was to see if it had any correlation with the location.

```

In [63]: EMS = df[(df['Reason'] == 'EMS') & (df['Hour'].isin(range(11,18)))]

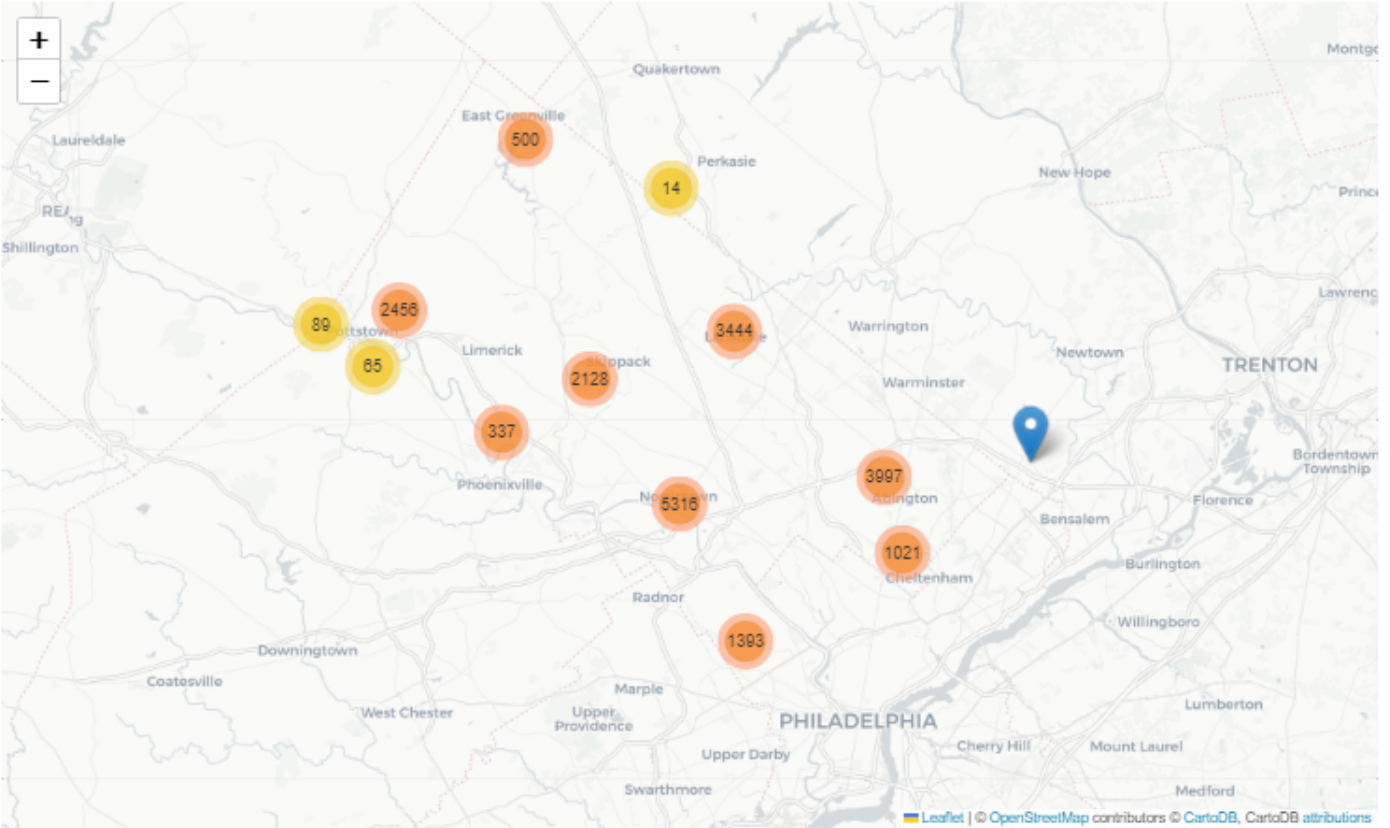
In [64]: # Create a base map
m_3 = folium.Map(location=[40.243821, -75.649367], tiles='cartodbpositron', zoom_start=8

# Add points to the map
mcd2 = MarkerCluster()
for idx, row in EMS.iterrows():
    if not math.isnan(row['lng']) and not math.isnan(row['lat']):
        mcd2.add_child(Marker([row['lat'], row['lng']]))
m_3.add_child(mcd2)

# Display the map
display.Image('EMS.png') #m_3

```

Out[64]:



Knowing that there are some concentrations in certain areas we could see several other indicators to understand how the place is or something like that.

Using the website for example: [https://www.city-data.com/county/Philadelphia\\_County-PA.html](https://www.city-data.com/county/Philadelphia_County-PA.html)

## Machine Learning Model

I focused on applying the K-Means method using only to the EMS data, since they represent almost 50% of the cases in the entire database.

```
In [65]: # Importing
from sklearn.cluster import KMeans
```

```
In [66]: # separating the data from EMS
df_EMS = df[df['Reason'] == 'EMS']
X = df_EMS[['lat', 'lng']]
X.head()
```

```
Out[66]:
```

	lat	lng
0	40.297876	-75.581294
1	40.258061	-75.264680
3	40.116153	-75.343513
4	40.251492	-75.603350
5	40.253473	-75.283245

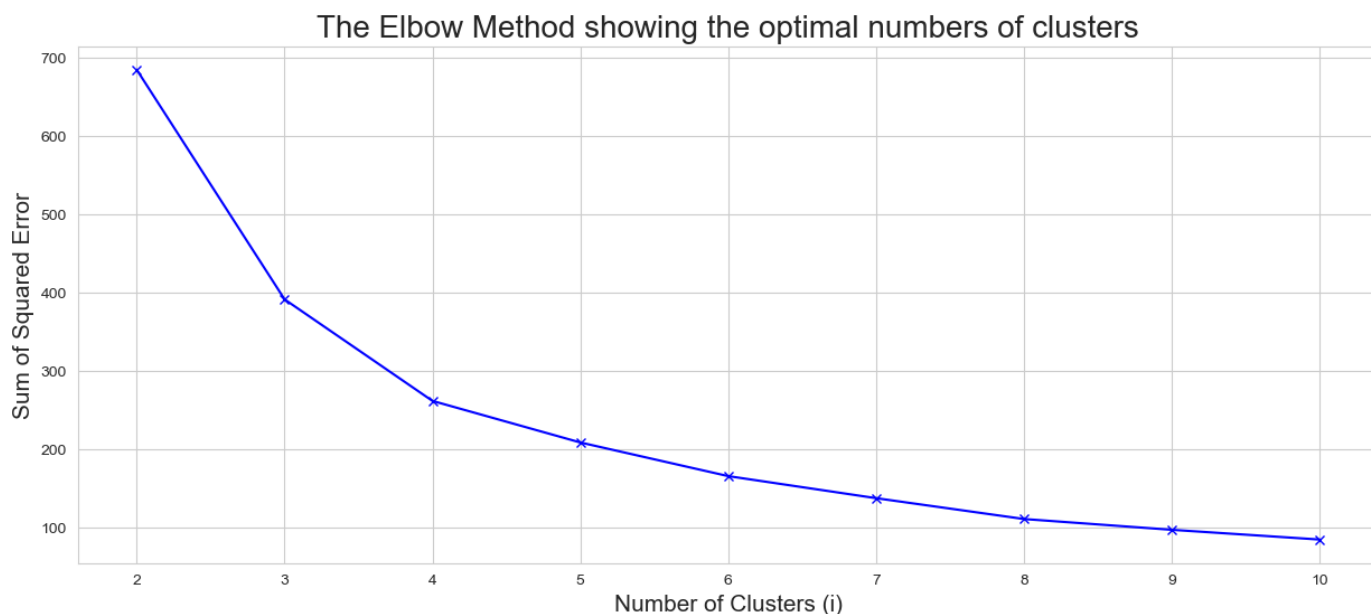
```
In [67]: # using elbow method to find the optimum cluster numbers in K-means
```



```
sse = []

K = range(2,11)
for i in K:
    kmeans = KMeans(n_clusters=i)
    kmeans.fit(X)
    sse.append(kmeans.inertia_)
```

```
In [68]: plt.figure(figsize=(15,6))
plt.plot(K, sse, 'bx-')
plt.xlabel('Number of Clusters (i)', fontsize = 15)
plt.ylabel('Sum of Squared Error', fontsize = 15)
plt.title('The Elbow Method showing the optimal numbers of clusters', fontsize = 20)
plt.show()
```



It was not very clear to identify the right 'K' using the elbow method. So, I used also the Silhouette score

```
In [69]: from sklearn.metrics import silhouette_score

X = df_EMS[['lat','lng']]
print("Clusters\tSilhoutte Score\n")
for n_cluster in range(2, 11):
    kmeans = KMeans(n_clusters=n_cluster).fit(X)
    label = kmeans.labels_
    sil_coeff = silhouette_score(X, label, metric='euclidean')
    print("k = {} \t--> \t{}".format(n_cluster, sil_coeff))
```

Clusters	Silhoutte Score
k = 2 -->	0.49432895139591215
k = 3 -->	0.4596900986644844
k = 4 -->	0.505550619809182
k = 5 -->	0.4800250732790006
k = 6 -->	0.4816194321919607
k = 7 -->	0.4984803500411124
k = 8 -->	0.5016000961943013
k = 9 -->	0.49510648074738955
k = 10 -->	0.46136006706515775

In the Elbow chart we can see that values above 3 would be a good fit, so in the Silhouette score we had that k = 4 gives the highest Silhouette Score

```
In [70]: kmeans = KMeans(n_clusters=4)
kmeans.fit(X)
```

```
sse.append(kmeans.inertia_)
```

```
In [71]: X['cluster_label'] = kmeans.fit_predict(X)
```

```
In [72]: X.head(3)
```

```
Out[72]:
```

	lat	lng	cluster_label
0	40.297876	-75.581294	2
1	40.258061	-75.264680	1
3	40.116153	-75.343513	0

```
In [75]: centers = pd.DataFrame(kmeans.cluster_centers_, columns=['lat', 'lng'])  
centers
```

```
Out[75]:
```

	lat	lng
0	40.098978	-75.333746
1	40.249289	-75.312572
2	40.260185	-75.570181
3	40.125895	-75.139225

```
In [76]: X['cluster_label'].value_counts()
```

```
Out[76]:
```

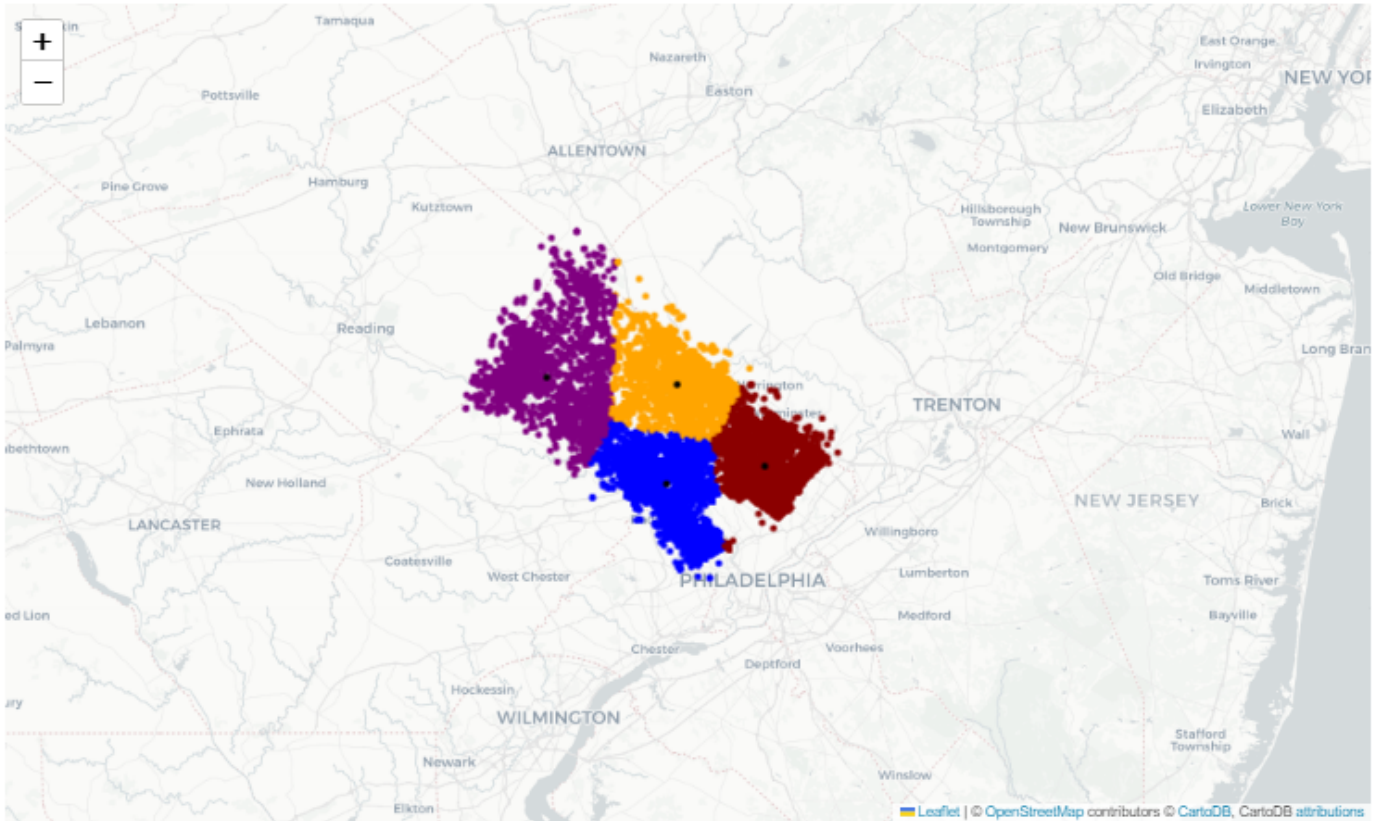
0	17080
3	12556
2	9753
1	9486

Name: cluster\_label, dtype: int64

```
In [80]: # Create a base map  
m_4 = folium.Map(location=[40.243821, -75.649367], tiles='cartodbpositron', zoom_start=8  
  
def color_producer(val):  
    if val == 0:  
        return 'blue'  
    elif val == 1:  
        return 'orange'  
    elif val == 2:  
        return 'purple'  
    else:  
        return 'darkred'  
  
# Add a bubble map to the base map  
for i in range(0, len(X)):  
    Circle(  
        location=[X.iloc[i]['lat'], X.iloc[i]['lng']],  
        radius=20,  
        color=color_producer(X.iloc[i]['cluster_label'])).add_to(m_4)  
  
for a in range(0, len(centers)):  
    Circle(  
        location=[centers.iloc[a]['lat'], centers.iloc[a]['lng']],  
        radius=30,  
        color='black').add_to(m_4)  
  
# Display the map  
display.Image('ems_kmeans.png') #m_4
```



Out[80]:



Here we can see how the model clustered the data. The black dots are the centroids, i.e. data points that represent the center of each cluster (the mean) and may not necessarily be a member of the data set.

**Therefore, considering only EMS, we must place the facilities to attend these cases close to the centroids to better respond to emergency calls, if there is only a budget for one point, it would be interesting that it was close to the centroid of Cluster 0 (blue), where we have a greater concentration of EMS cases.**

---

## Thanks!

If you've found some english error or something that you want to share, please tell me